

Declarative programming over eventually consistent data stores

CMPS290S, Fall 2018

Lindsey Kuper

November 9, 2018

main idea

KC Sivaramakrishnan



Hey, what if we took
Burckhardt and co.'s RDT
specification framework¹...
...and *turned it into a
programming language?!*

¹ you know, the *vis* relation and all that stuff,
as seen in Gotsman et al.'s POPL '14 paper,
Burckhardt's book, the Viotti and Vukolić
survey, and probably a bunch of other places

main idea

In particular, what if we turned it into a *contract language* for specifying the *consistency semantics* of data stores and operations on them?

KC Sivaramakrishnan



main idea

KC Sivaramakrishnan



In particular, what if we turned it into a *contract language* for specifying the *consistency semantics* of data stores and operations on them?

...and implemented it in Haskell?



main idea

KC Sivaramakrishnan



In particular, what if we turned it into a *contract language* for specifying the *consistency semantics* of data stores and operations on them?

...and implemented it in Haskell?



...and stuck it on top of Cassandra?



main idea

KC Sivaramakrishnan



In particular, what if we turned it into a *contract language* for specifying the *consistency semantics* of data stores and operations on them?

...and implemented it in Haskell?



...and stuck it on top of Cassandra?



yeah
sounds good
let's do that

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```



operations

every op gets a *history* of known updates to the object as its first argument

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```



operations

every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```



operations

every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```



operations

every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations

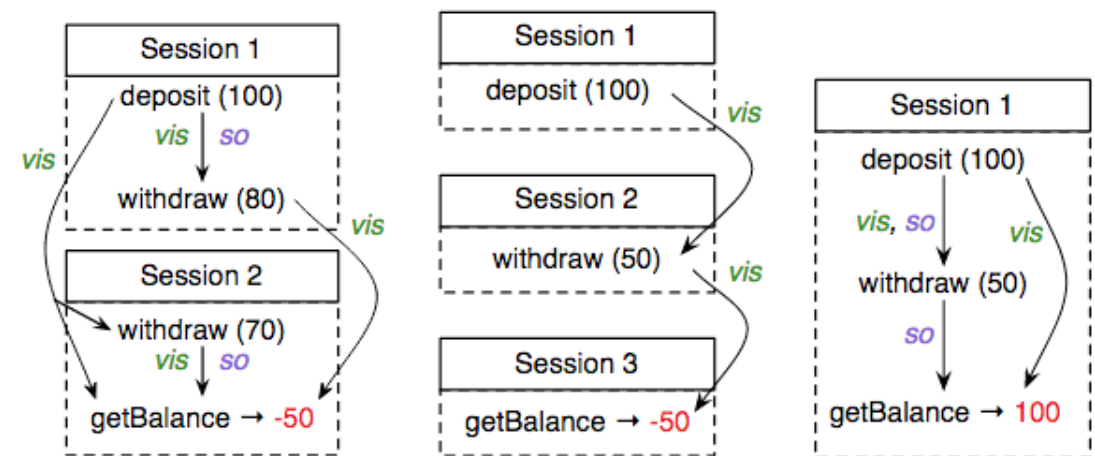
every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations


every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



violation of *integrity invariant* ($\text{balance} \geq 0$)
fixable w/ strong consistency
(or token system, as in the CISE paper)

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations


every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations

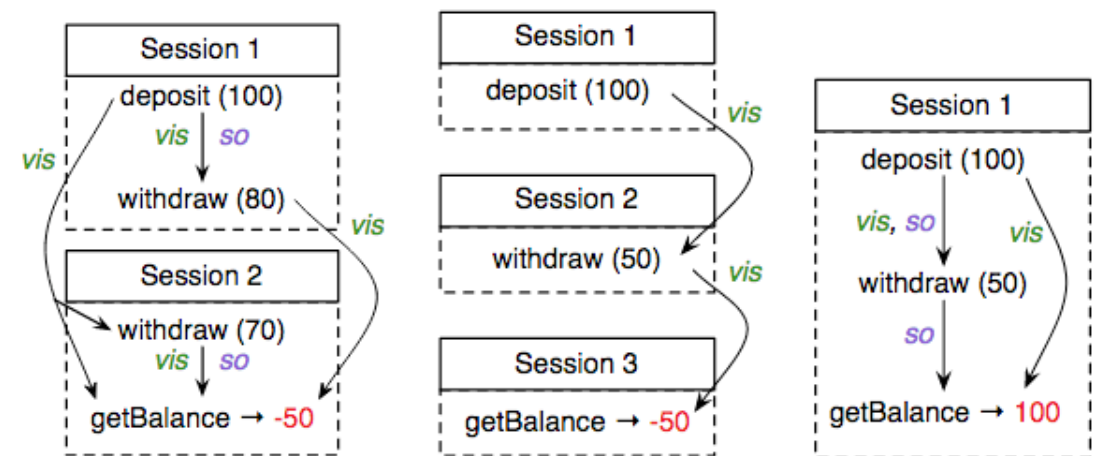
every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



violation of *causality*
fixable w/ causal consistency

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations

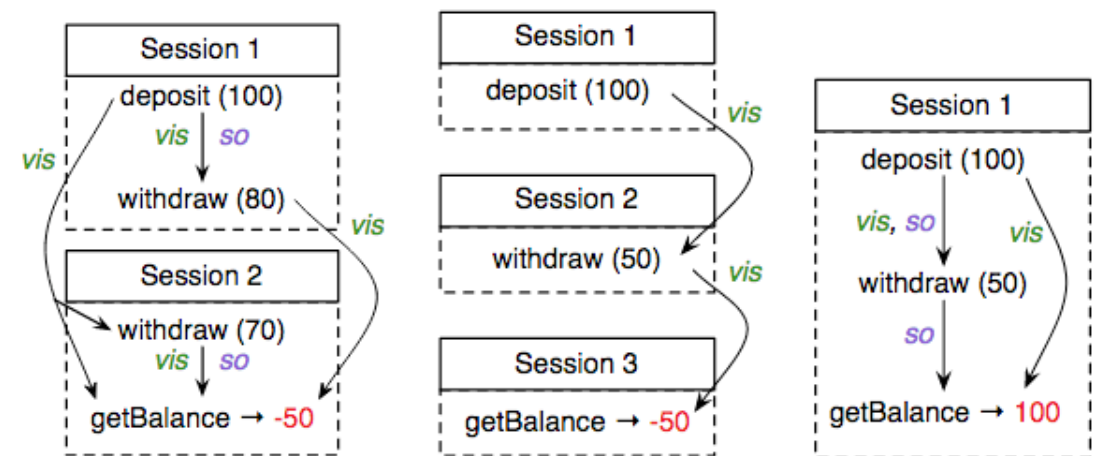
every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations


every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



violation of RYW session guarantee
(and also causality)
fixable w/ causal consistency

the ubiquitous bank account example

```
data Acc = Deposit Int | Withdraw Int | GetBal
```

```
getBalance :: [Acc] → () → (Int, Maybe Acc)
getBalance hist _ =
  let res = sum [x | Deposit x ← hist] -
            sum [x | Withdraw x ← hist]
  in (res, Nothing)
```

```
deposit :: [Acc] → Int → ((), Maybe Acc)
deposit hist amt = ((), Just $ Deposit amt)
```

```
withdraw :: [Acc] → Int → (Bool, Maybe Acc)
withdraw hist v =
  if sell $ getBalance hist () ≥ v
  then (True, Just $ Withdraw v)
  else (False, Nothing)
```

operations

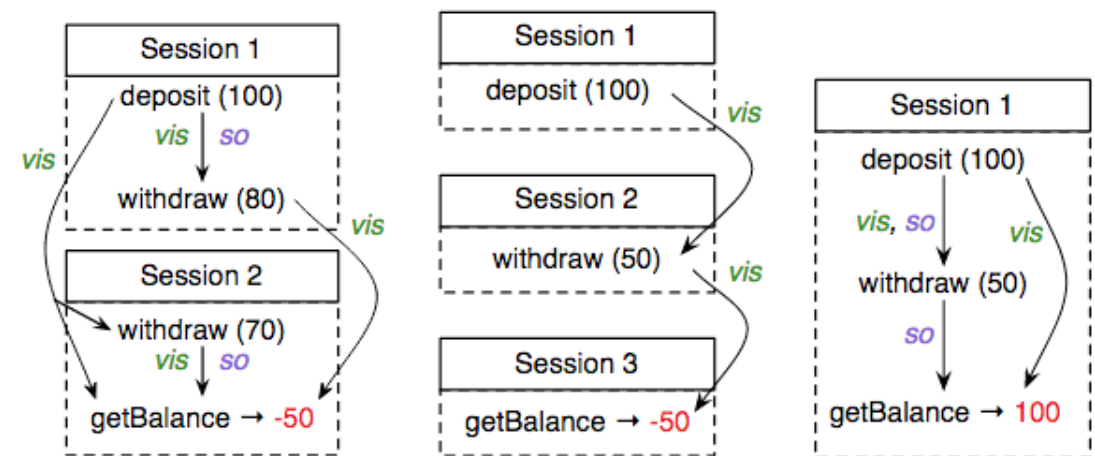
every op gets a *history* of known updates to the object as its first argument

Q for class: what consistency level do each of these three ops need?

summarization: every now and then (how often?) the history gets boiled down

Q for class: how is summarization done for this particular RDT?

let's play “name that anomaly”



the contract language

$$\begin{aligned}\psi &\in \text{Contract} & ::= & \forall(x : \tau).\psi \mid \forall x.\psi \mid \pi \\ \tau &\in \text{EffType} & ::= & \text{Op} \mid \tau \vee \tau \\ \pi &\in \text{Prop} & ::= & \text{true} \mid R(x, y) \mid \pi \vee \pi \\ & & & \mid \pi \wedge \pi \mid \pi \Rightarrow \pi \\ R &\in \text{Relation} & ::= & \text{vis} \mid \text{so} \mid \text{sameobj} \mid = \\ & & & \mid R \cup R \mid R \cap R \mid R^+\end{aligned}$$
$$x, y, \hat{\eta} \in \text{EffVar} \qquad \text{Op} \in \text{OpName}$$

the contract language

a contract is a first-order logic formula
universal quantification over EffVars allowed

$\psi \in \text{Contract} ::= \forall(x : \tau).\psi \mid \forall x.\psi \mid \pi$
 $\tau \in \text{EffType} ::= \text{Op} \mid \tau \vee \tau$
 $\pi \in \text{Prop} ::= \text{true} \mid R(x, y) \mid \pi \vee \pi$
 $\quad \mid \pi \wedge \pi \mid \pi \Rightarrow \pi$
 $R \in \text{Relation} ::= \text{vis} \mid \text{so} \mid \text{sameobj} \mid =$
 $\quad \mid R \cup R \mid R \cap R \mid R^+$

$x, y, \hat{\eta} \in \text{EffVar}$

$\text{Op} \in \text{OpName}$

the contract language

$\psi \in \text{Contract} ::= \forall(x : \tau).\psi \mid \forall x.\psi \mid \pi$ a contract is a first-order logic formula
universal quantification over EffVars allowed

$\tau \in \text{EffType} ::= \text{Op} \mid \tau \vee \tau$

$\pi \in \text{Prop} ::= \text{true} \mid R(x, y) \mid \pi \vee \pi$
 $\mid \pi \wedge \pi \mid \pi \Rightarrow \pi$ hey look it's our old friends from Burckhardt et al.

$R \in \text{Relation} ::= \text{vis} \mid \text{so} \mid \text{sameobj} \mid =$
 $\mid R \cup R \mid R \cap R \mid R^+$

$x, y, \hat{\eta} \in \text{EffVar} \quad \text{Op} \in \text{OpName}$

the contract language

$\psi \in \text{Contract} ::= \forall(x : \tau).\psi \mid \forall x.\psi \mid \pi$ a contract is a first-order logic formula
universal quantification over EffVars allowed

$\tau \in \text{EffType} ::= \text{Op} \mid \tau \vee \tau$

$\pi \in \text{Prop} ::= \text{true} \mid R(x, y) \mid \pi \vee \pi$
 $\mid \pi \wedge \pi \mid \pi \Rightarrow \pi$ hey look it's our old friends from Burckhardt et al.

$R \in \text{Relation} ::= \text{vis} \mid \text{so} \mid \text{sameobj} \mid =$
 $\mid R \cup R \mid R \cap R \mid R^+$

$x, y, \hat{\eta} \in \text{EffVar}$

$\text{Op} \in \text{OpName}$

interesting aside: this is actually a lie
the contract language can't express transitive closure!
(Q for class: why not?)

the contract language

$\psi \in \text{Contract} ::= \forall(x : \tau).\psi \mid \forall x.\psi \mid \pi$ a contract is a first-order logic formula
universal quantification over EffVars allowed
 $\tau \in \text{EffType} ::= \text{Op} \mid \tau \vee \tau$
 $\pi \in \text{Prop} ::= \text{true} \mid R(x, y) \mid \pi \vee \pi$
 $\quad \quad \quad \mid \pi \wedge \pi \mid \pi \Rightarrow \pi$ hey look it's our old friends from Burckhardt et al.
 $R \in \text{Relation} ::= \text{vis} \mid \text{so} \mid \text{sameobj} \mid =$
 $\quad \quad \quad \mid R \cup R \mid R \cap R \mid R^+$

$x, y, \hat{\eta} \in \text{EffVar}$

$\text{Op} \in \text{OpName}$

interesting aside: this is actually a lie
the contract language can't express transitive closure!
(Q for class: why not?)

some example contracts for operations

($\hat{\eta}$ is the current operation/effect)

for withdraw: $\forall(a : \text{withdraw}).$
 $\text{sameobj}(a, \hat{\eta}) \Rightarrow a = \hat{\eta} \vee \text{vis}(a, \hat{\eta}) \vee \text{vis}(\hat{\eta}, a)$

for getBalance: $\forall(a : \text{deposit}), (b : \text{withdraw}), (c : \text{deposit} \vee \text{withdraw}).$
 $(\text{vis}(a, b) \wedge \text{vis}(b, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta}))$
 $\wedge ((\text{so} \cap \text{sameobj})(c, \hat{\eta}) \Rightarrow \text{vis}(c, \hat{\eta}))$

the contract language

$\psi \in \text{Contract} ::= \forall(x : \tau).\psi \mid \forall x.\psi \mid \pi$ a contract is a first-order logic formula
 universal quantification over EffVars allowed
 $\tau \in \text{EffType} ::= \text{Op} \mid \tau \vee \tau$
 $\pi \in \text{Prop} ::= \text{true} \mid R(x, y) \mid \pi \vee \pi$
 $\mid \pi \wedge \pi \mid \pi \Rightarrow \pi$ hey look it's our old friends from Burckhardt et al.
 $R \in \text{Relation} ::= \text{vis} \mid \text{so} \mid \text{sameobj} \mid =$
 $\mid R \cup R \mid R \cap R \mid R^+$

$x, y, \hat{\eta} \in \text{EffVar}$

$\text{Op} \in \text{OpName}$

interesting aside: this is actually a lie
 the contract language can't express transitive closure!
 (Q for class: why not?)

some example contracts for operations

($\hat{\eta}$ is the current operation/effect)

for withdraw: $\forall(a : \text{withdraw}).$
 $\text{sameobj}(a, \hat{\eta}) \Rightarrow a = \hat{\eta} \vee \text{vis}(a, \hat{\eta}) \vee \text{vis}(\hat{\eta}, a)$

for getBalance: $\forall(a : \text{deposit}), (b : \text{withdraw}), (c : \text{deposit} \vee \text{withdraw}).$
 $(\text{vis}(a, b) \wedge \text{vis}(b, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta}))$
 $\wedge ((\text{so} \cap \text{sameobj})(c, \hat{\eta}) \Rightarrow \text{vis}(c, \hat{\eta}))$

Q for class: what's the contract for deposit? why?

stores have contracts, too

eventually consistent store: $\psi_{ec} = \forall a, b. \text{hbo}(a, b) \wedge \text{vis}(b, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta})$

causally consistent store: $\psi_{cc} = \forall a. \text{hbo}(a, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta})$

strongly consistent store: $\psi_{sc} = \forall a. \text{sameobj}(a, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta}) \vee \text{vis}(\hat{\eta}, a) \vee a = \hat{\eta}$

stores have contracts, too

eventually consistent store: $\psi_{ec} = \forall a, b. \text{hbo}(a, b) \wedge \text{vis}(b, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta})$

Q for class:

as definitions of EC go, this one is pretty strong!

what executions does it admit that causal consistency *doesn't*?

causally consistent store: $\psi_{cc} = \forall a. \text{hbo}(a, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta})$

strongly consistent store: $\psi_{sc} = \forall a. \text{sameobj}(a, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta}) \vee \text{vis}(\hat{\eta}, a) \vee a = \hat{\eta}$

stores have contracts, too

eventually consistent store: $\psi_{ec} = \forall a, b. \text{hbo}(a, b) \wedge \text{vis}(b, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta})$

Q for class:

as definitions of EC go, this one is pretty strong!

what executions does it admit that causal consistency *doesn't*?

causally consistent store: $\psi_{cc} = \forall a. \text{hbo}(a, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta})$

strongly consistent store: $\psi_{sc} = \forall a. \text{sameobj}(a, \hat{\eta}) \Rightarrow \text{vis}(a, \hat{\eta}) \vee \text{vis}(\hat{\eta}, a) \vee a = \hat{\eta}$

Qs for class:

does this one remind you of anything from another paper we read recently?

does it correspond to anything in the Viotti and Vukolić zoo?

contract classification

$$\begin{array}{cc} \frac{\psi \leq \psi_{sc}}{\text{WellFormed}(\psi)} & \frac{\psi \leq \psi_{ec}}{\text{EventuallyConsistent}(\psi)} \\ \frac{\psi \not\leq \psi_{ec} \quad \psi \leq \psi_{cc}}{\text{CausallyConsistent}(\psi)} & \frac{\psi \not\leq \psi_{cc} \quad \psi \leq \psi_{sc}}{\text{StronglyConsistent}(\psi)} \end{array}$$

note: this really is “classification”,
in the sense that each contract gets classified to *exactly one* consistency level
(i.e., if it’s CausallyConsistent then it is *not also* EventuallyConsistent)

could also swap in a different set of store consistency levels,
e.g, the four session guarantees,
or even various transaction isolation levels (see paper)

contract classification

$$\begin{array}{cc} \frac{\psi \leq \psi_{sc}}{\text{WellFormed}(\psi)} & \frac{\psi \leq \psi_{ec}}{\text{EventuallyConsistent}(\psi)} \\ \frac{\psi \not\leq \psi_{ec} \quad \psi \leq \psi_{cc}}{\text{CausallyConsistent}(\psi)} & \frac{\psi \not\leq \psi_{cc} \quad \psi \leq \psi_{sc}}{\text{StronglyConsistent}(\psi)} \end{array}$$

note: this really is “classification”,
in the sense that each contract gets classified to *exactly one* consistency level
(i.e., if it’s CausallyConsistent then it is *not also* EventuallyConsistent)

could also swap in a different set of store consistency levels,
e.g, the four session guarantees,
or even various transaction isolation levels (see paper)

Qs for class:

how would you classify the `getBalance` contract?

how about the `withdraw` contract?

would Owen approve of the lattice in figure 7?

soundness of contract enforcement

Theorem 4 (Soundness of Contract Enforcement). *Let ψ be a well-formed contract of a replicated data type operation op , and let τ denote the consistency class of ψ as determined by the contract classification scheme. For all well-formed execution states E, E' such that $E, \langle op, \tau \rangle; \sigma \parallel \Sigma \xrightarrow{\eta} E', \sigma \parallel \Sigma$, if $E' \models \psi_\tau[\eta/\hat{\eta}]$, then $E' \models \psi[\eta/\hat{\eta}]$*

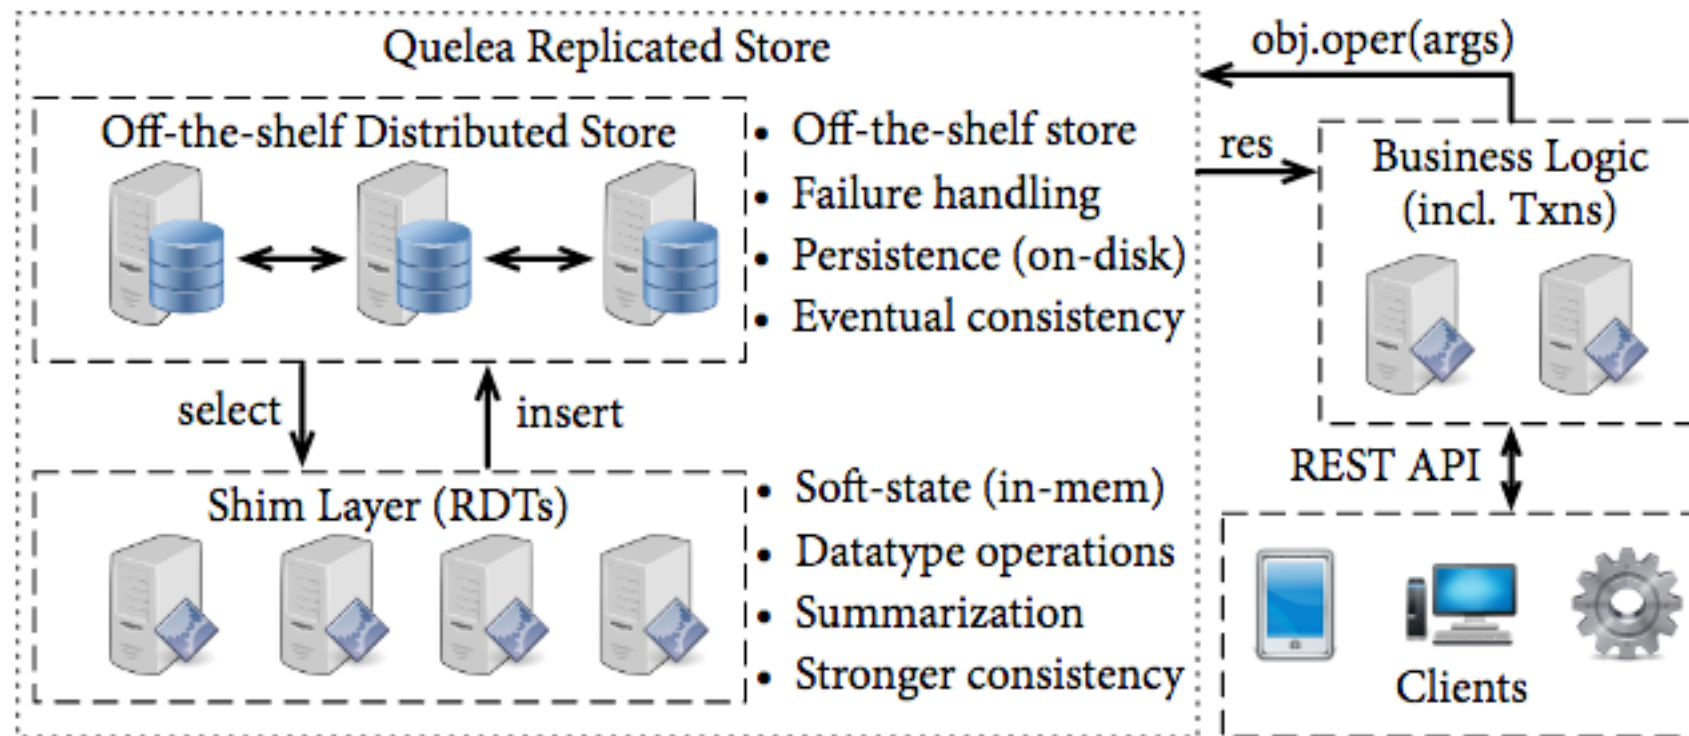

simpler:

“If you take a step and execute an operation whose contract has been classified in a certain way, and the resulting execution state satisfies that consistency classification, then the execution state also satisfies the operation’s original contract.”

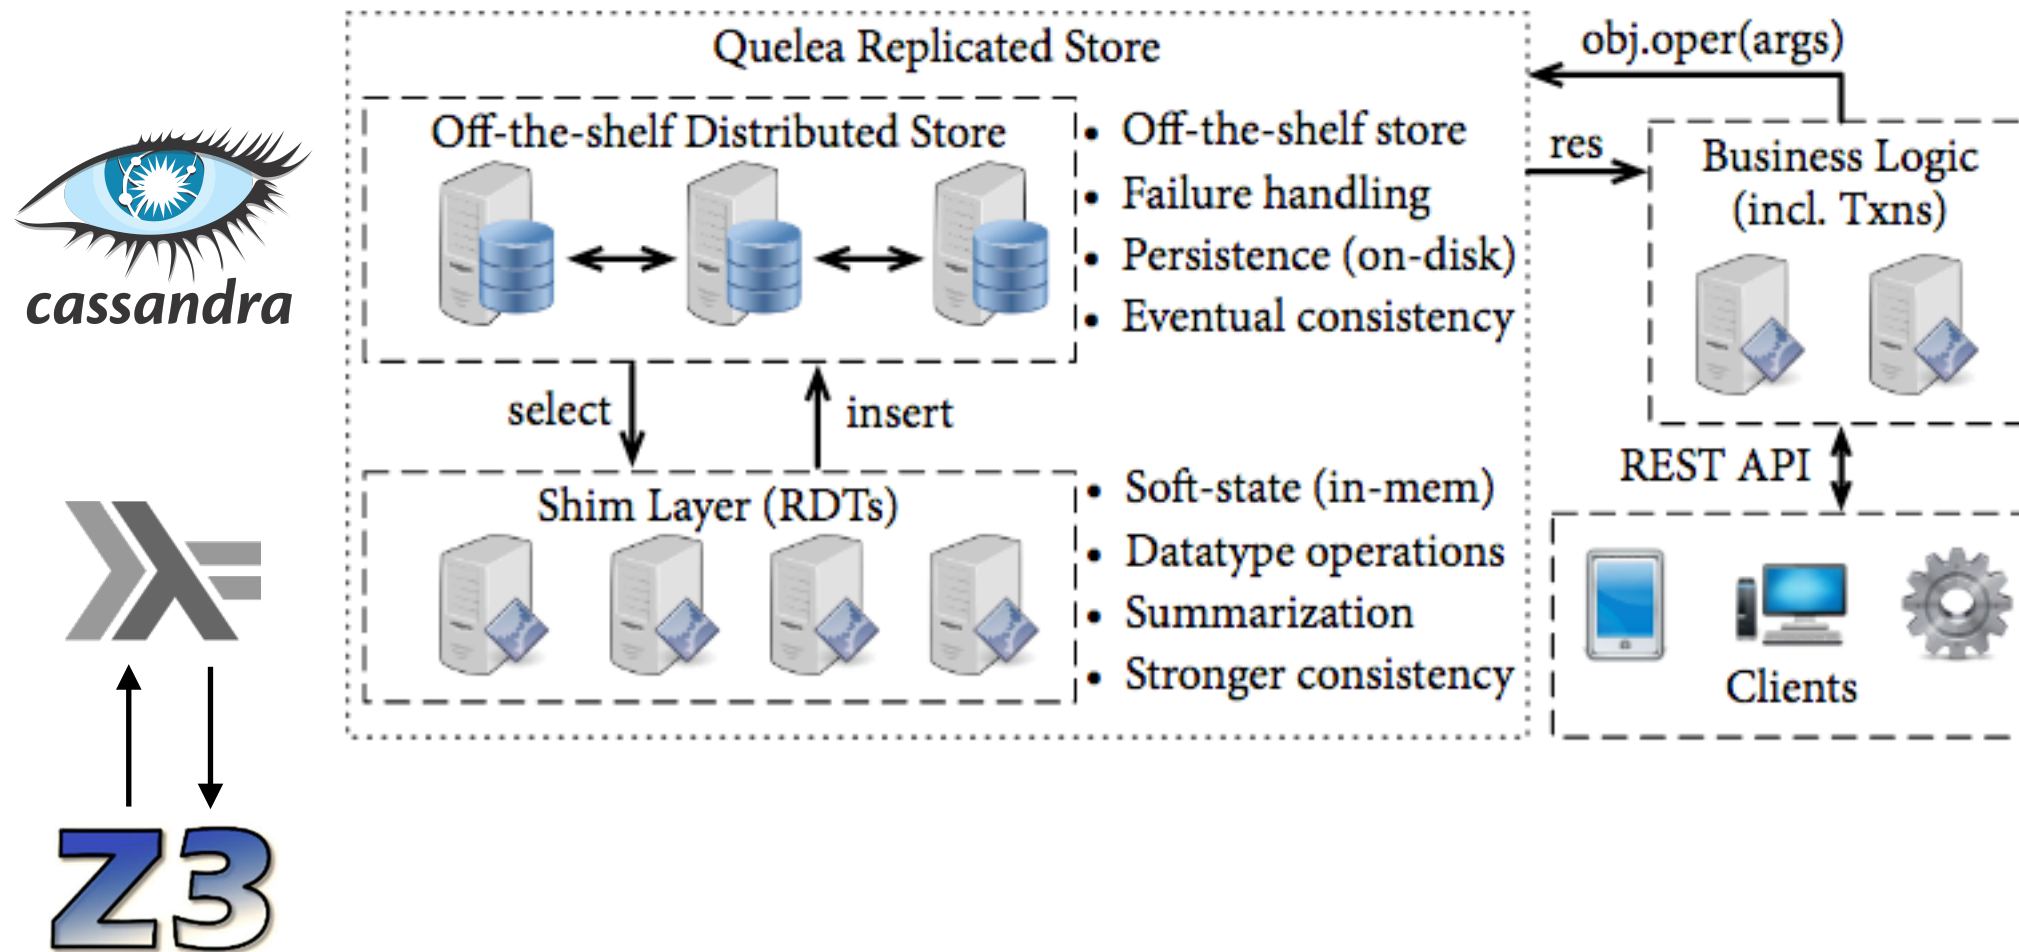
or:

“If an operation’s contract is in a particular store consistency class, and running the operation results in a state satisfying that store consistency, then the operation’s contract is satisfied.”

implementation




implementation



Qs for class:

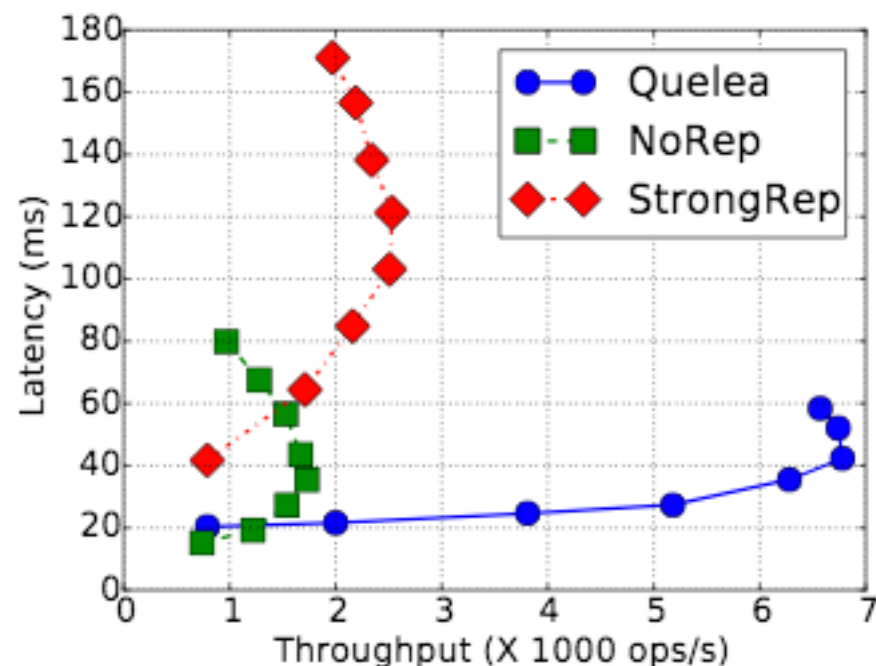
**how can one add causal consistency atop Cassandra's quorum consistency?
(could Holt et al.'s IPA have done the same thing?)**

evaluation



nice result:

Quelea does a bit better, throughput-wise, than doing everything at causal consistency (which isn't even "correct")!



Quelea has lower throughput *and* latency than *not replicating*! (why?)

my questions

how often should one summarize?

individual operation contracts can be really fine-grained, only to be classified into more coarse-grained store contract levels. how about a more granular way of interacting with the store?

what do they wish they could've had in the contract language that they couldn't have because of the limitations of SMT?