

Modular Software Upgrades For Distributed Systems

Authors:

Sameer Ajmani
Barbara Liskov
Liuba Shrira

Presentation By:

Aldrin Montana

Requirements

1. Modularity
2. Generality
3. Transform persistent state (Schema Evolution?)
4. Automatic Deployment
5. Controlled Deployment
6. Continuous Service (Zero downtime)

A hand-drawn illustration of a whiteboard with the text "WE UNDERSTAND YOUR NEEDS" written in black marker. The text is written in a casual, slightly slanted font. The word "NEEDS" is underlined. A hand is visible at the bottom right, holding a black marker and finishing the underline.

WE UNDERSTAND
YOUR NEEDS

Contributions

1. First complete solution for automatic and controlled upgrades in distributed systems
2. A way to understand and specify mixed mode

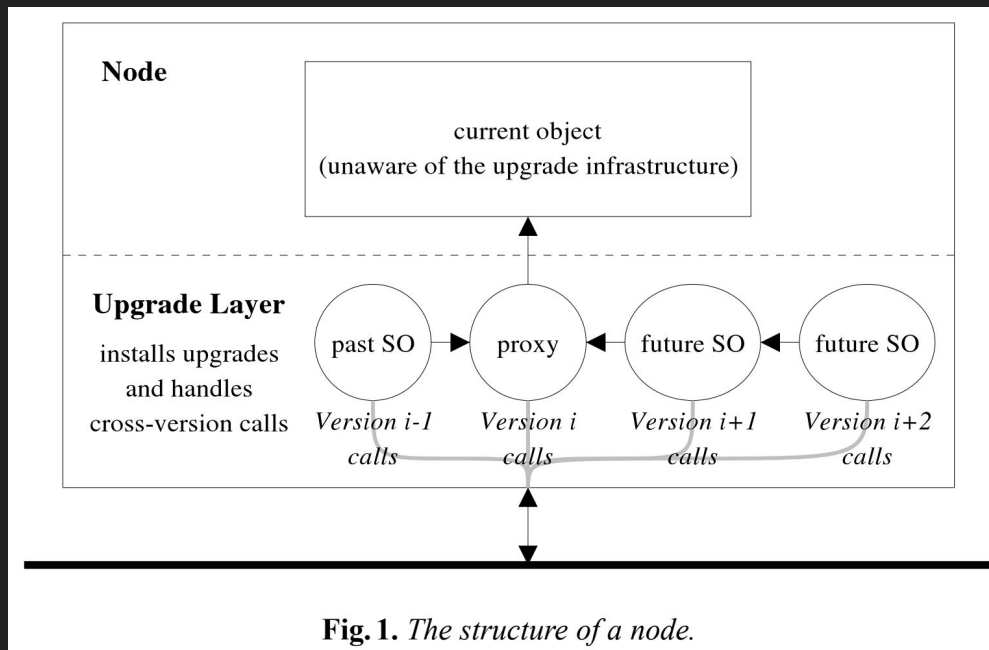


Data Model

- Distributed system is a collection of objects
- An object has:
 - Identity
 - Type (defines behavior)
 - State (a portion may be persistent)
 - Object reinitializes itself from persistent state
- Objects call each other's methods


Data Model

An object is instance of a class that defines how it implements its type



Terms

- Simulation Object
- Delegate Object
- Invariant
- Type
- Event
- Call



Type Requirement

$O_1.m(args); O_1.m(args);$ [version 2 introduced at server];
 $O_1.m(args); O_2.p(args);$ [server upgrades from 1 to 2];
 $O_1.m(args); O_2.p(args);$ [version 1 retired];
 $O_2.p(args); O_2.p(args);$

Type Requirement

- The class implementing a future SO must implement the new type
- The class implementing a past SO must implement the old type
- Effects of interleaving need to be defined

$O_1.m(args); O_1.m(args);$ [version 2 introduced at server];
 $O_1.m(args); O_2.p(args);$ [server upgrades from 1 to 2];
 $O_1.m(args); O_2.p(args);$ [version 1 retired];
 $O_2.p(args); O_2.p(args);$

Sequence and Disallow

- Sequence Requirement
 - An event is a call, upgrade, or introduction of a version
 - Calls to a **current object or SO** must reflect **effects of calls to other objects**
- Disallow Constraint
 - Calls to current object cannot be disallowed
 - Current object provides “real behavior”

Specifying Upgrades

- Three Parts:

Specifying Upgrades

- Three Parts:
 - a. Invariant
 - b. Mapping function
 - c. Shadow methods

Specifying Upgrades - Invariant

- Corresponding set of old and new object state
- Must be **total**
 - a. One-to-One?
 - b. For each legal state O_{new} of T_{new} , there is some legal state, O_{old} of T_{old}
 - c. And vice versa

Specifying Upgrades - Mapping Function

- Initial state for O_{new} given state of O_{old}
- Must be total and establish the invariant

State of new file system given old file system?

Specifying Upgrades - Shadow Methods

- Corresponding set of old and new methods
 - a. Preconditions must hold
 - b. Invariant must be preserved



Examples - ColorSet with direct simulation

- `SO.insertColor(1, red)`
- `O.delete(1)`
- `O.insert(1)`
- `SO.getColor(1)`

Examples - ColorSet with direct simulation

- `SO.insertColor(1, red)`
- `O.delete(1)`
- `O.insert(1)`
- `SO.getColor(1)`

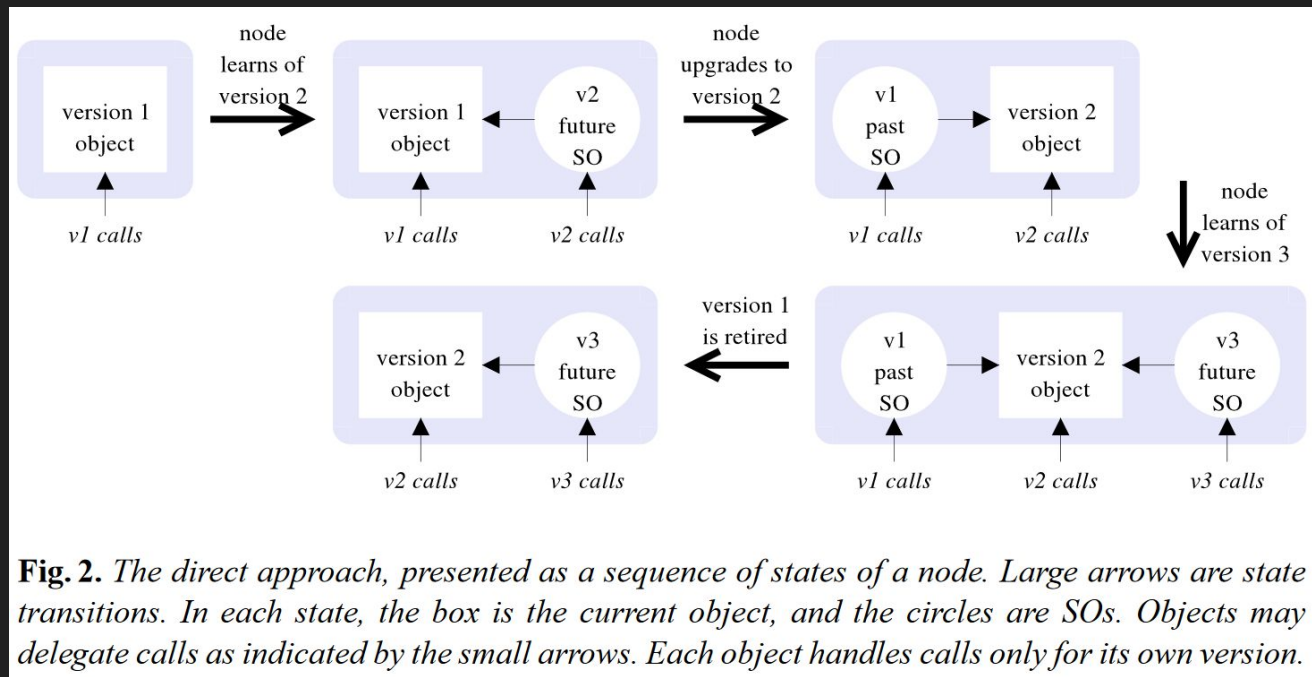


Fig. 2. The direct approach, presented as a sequence of states of a node. Large arrows are state transitions. In each state, the box is the current object, and the circles are SOs. Objects may delegate calls as indicated by the small arrows. Each object handles calls only for its own version.

Examples - interceptor simulation

- `SO.insertColor(1, red)`
- `O.delete(1)`
- `O.insert(1)`
- `SO.getColor(1)`

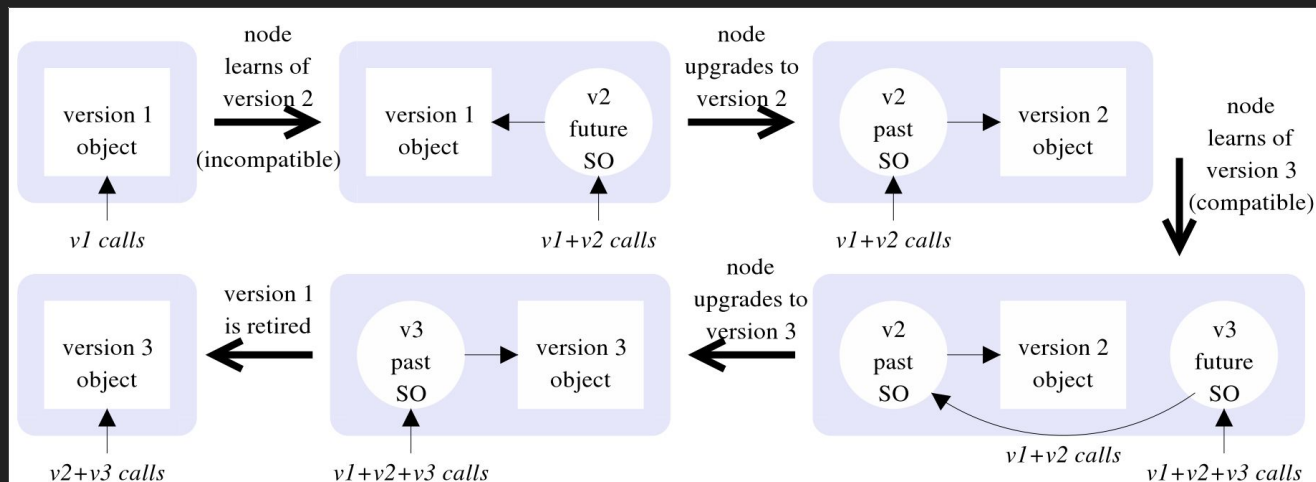


Fig. 3. The interceptor approach, presented as a sequence of states of a node. Large arrows are state transitions. In each state, the box is the current object, and the circles are SOs. Objects may delegate calls as indicated by the small arrows. One object in each state intercepts all calls.

Transform and Scheduling Functions

- TF must implement identity mapping
 - post-TF of past SO \longleftrightarrow pre-TF of old object
 - post-TF of new object \longleftrightarrow pre-TF of future SO
 - Must be restartable (idempotent?)
- SF runs on node itself
 - Access to UDB (central knowledge) and LDB (local knowledge)
 - Strategies: Rolling upgrade, Big flip, Fast reboot

Evaluation - microbenchmarks

The overhead imposed by the upgrade layer
when no upgrades are happening

	Null RPC (loopback)			Null RPC (crossover)			100MB TCP transfer		
	5%	50%	95%	5%	50%	95%	5%	50%	95%
Baseline	50 μ s	51 μ s	53 μ s	247 μ s	382 μ s	769 μ s	896ms	896ms	923ms
TESLA	128 μ s	139 μ s	154 μ s	371 μ s	382 μ s	782 μ s	896ms	898ms	919ms
Upstart	192 μ s	206 μ s	223 μ s	245 μ s	388 μ s	819 μ s	897ms	908ms	936ms

Table 1. Microbenchmark results ($N=100000$ for Null RPC, $N=100$ for TCP). For each experiment, the 5th, 50th (median), and 95th percentile latencies are given.

Discuss